

Fundamentals of Information Systems

Python Programming (for Data Science)

Master's Degree in Data Science

Gabriele Tolomei

gtolomei@math.unipd.it

University of Padua, Italy

2018/2019

October 18, 2018

Lecture 3: Python's Built-in Data Types (1)

Data Type Hierarchy

- Python's built-in data types can be grouped into several classes.
- We use the same hierarchy scheme used in the [official Python documentation](#), which defines the following classes:
 - **numeric, sequences, sets and mappings** (and a few more not discussed further here).
- A special mention goes to two particular data types: **bool** and **NoneType**.

Booleans

Type `bool` (*immutable*)

- It encapsulates the two boolean values which are written as `True` and `False`.
- Comparisons and other conditional expressions evaluate to either `True` or `False`.
- Boolean values are combined with the `and` and `or` keywords.

```
In [1]: type(True)
```

```
Out[1]: bool
```

Boolean Operations: `or`, `and`, `not`

- Ordered by ascending priority:

Operation	Result
<code>x or y</code>	if <code>x</code> is false, then <code>y</code> , else <code>x</code>
<code>x and y</code>	if <code>x</code> is false, then <code>x</code> , else <code>y</code>
<code>not x</code>	if <code>x</code> is false, then <code>True</code> , else <code>False</code>

In [2]: **False or True**

Out[2]: True

In [3]: `True and True`

Out[3]: True

Comparisons

- There are **eight** comparison operations in Python.
- They all have the same priority (which is higher than that of the Boolean operations).
- Comparisons can be chained arbitrarily; for example, `x < y <= z` is equivalent to `x < y and y <= z`, except that `y` is evaluated **only once** (but in both cases `z` is not evaluated at all when `x < y` is found to be **False**).

Table of Comparisons

Operation	Meaning
<	strictly less than
<=	less than or equal
>	strictly greater than
>=	greater than or equal
==	equal
!=	not equal
is	object identity
is not	negated object identity

A Quick Note on the `is` Operator

- It is used to compare the **identity** of two objects.
- The **identity** of an object can be found with the `id()` built-in function.
- `id()` takes as input a Python object and returns an integer representing the identity of *that* object.
- In the standard CPython implementation, this integer corresponds to the object's location in memory (in other implementations/platforms this might be different).

`is` VS. `==`

- `is` is used to test for **identity** of two objects by means of the `id()` function.
- `==` is used to test for the **value** of two objects.
- In other words, if you have 2 objects `x` and `y` the statement below

`x is y`

corresponds to the following:

`id(x) == id(y)`

In [4]:

```
# Using the 'is' operator in combination with immutable objects (e.g., integers)
x = 42
y = x
print("id(x) = {}".format(id(x)))
print("id(42) = {}".format(id(42)))
print("id(y) = {}".format(id(y)))
print("Q: The identity of x is the same of that of y? A: {}".format(x is y)) # id(x) == id(y)
# Modifying x (immutable) means creating a new integer object and assign it to x
x += 1
print("id(x) = {}".format(id(x)))
print("id(43) = {}".format(id(43)))
print("id(42) = {}".format(id(42)))
print("id(y) = {}".format(id(y)))
print("Q: The identity of x is the same of that of y? A: {}".format(x is y)) # id(x) == id(y)
```

id(x) = 4460251760

id(42) = 4460251760

id(y) = 4460251760

Q: The identity of x is the same of that of y? A: True

id(x) = 4460251792

id(43) = 4460251792

id(42) = 4460251760

id(y) = 4460251760

Q: The identity of x is the same of that of y? A: False

In [5]:

```
# Using the 'is' operator in combination with mutable objects (e.g., lists)
x = [1, 2, 3]
y = x
print("id(x) = {}".format(id(x)))
print("id(y) = {}".format(id(y)))
print("Q: The identity of x is the same of that of y? A: {}".format(x is y)) # id(x) == id(y)
# Let's modify x
x.append(4)
print("id(x) = {}".format(id(x)))
print("id(y) = {}".format(id(y)))
print("Q: The identity of x is the same of that of y? A: {}".format(x is y)) # id(x) == id(y)
```

id(x) = 4502699464

id(y) = 4502699464

Q: The identity of x is the same of that of y? A: True

id(x) = 4502699464

id(y) = 4502699464

Q: The identity of x is the same of that of y? A: True

In [6]:

```
# Unexpected behaviors which might cause you some problems...
x = 42
y = 42
print("Q: The identity of x is the same of that of y? A: {}".format(x is y)) # id(x) == id(y)
print("Q: The value of x is the same of that of y? A: {}".format(x == y))
x = 257
y = 257
print("Q: The identity of x is the same of that of y? A: {}".format(x is y)) # id(x) == id(y)
print("Q: The value of x is the same of that of y? A: {}".format(x == y))
# This odd behavior depends on the fact that CPython implements
# integers in the range (-5, 256) at fixed memory locations. As such, any named variable
# referencing one of those integers will always have the same memory address.
# On the other hand, integers outside that range might be possibly allocated at different
# memory addresses and therefore they have different identities even though the same value!
# Long story short, if you want to test for equality DO USE '=='
```

Q: The identity of x is the same of that of y? A: True

Q: The value of x is the same of that of y? A: True

Q: The identity of x is the same of that of y? A: False

Q: The value of x is the same of that of y? A: True

In [7]:

```
# When you work with mutable objects you will always face the following behavior
x = [1, 2, 3]
y = [1, 2, 3] # Note that here we are assigning a 'new' object to y
print("Q: The identity of x is the same of that of y? A: {}".format(x is y)) # id(x) == id(y)
print("Q: The value of x is the same of that of y? A: {}".format(x == y))
```

Q: The identity of x is the same of that of y? A: False

Q: The value of x is the same of that of y? A: True

Non-zero Interpretation

- Almost all built-in Python types (and any class defining the `__nonzero__` method) have a `True` or `False` interpretation in an `if` statement

In [8]:

```
1 x = [1, 2, 3] # define a list with 3 elements
2 if x:
3     print('The list contains something!')
4
5 y = [] # define an empty list
6 if not y:
7     print('The list is empty!')
```

The list contains something!

The list is empty!

True- or Falseness

- Most objects in Python have a notion of true- or falseness.
- For example, empty sequences like lists, dicts, tuples, etc. (more on those types later on) are treated as **False** if used in control flow (see the empty list **y** above).
- You can see exactly what boolean value an object coerces to by invoking **bool** on it.

```
In [9]: bool([], bool([1, 2, 3]))
```

```
Out[9]: (False, True)
```

```
In [10]: bool('Hello World!'), bool('')
```

```
Out[10]: (True, False)
```

```
In [11]: bool(0), bool(1)
```

```
Out[11]: (False, True)
```

None

Type `NoneType` (*immutable*) and `None` instance

- `None` is the Python null value type.
- Actually, it is the unique available *instance* of `NoneType` object.
- If a function does not explicitly return a value, it implicitly returns `None`.
- `None` is also a common default value for *optional* function arguments.

In [12]:

```
a = None  
a is None
```

Out[12]: True

```
In [13]: b = 42  
b is None
```

```
Out[13]: False
```

In [14]:

```
1 # z is an optional input argument of the following function
2 def add_and_possibly_multiply(x, y, z=None):
3
4     result = x + y # sum the first two positional input arguments
5
6     if z is not None: # multiply the current result by z iff z is not None
7         result *= z
8
9     return result # finally, return the result
```

Numerics

Numeric Types: `int`, `float`, `complex` (*immutables*)

- The primary Python types for numbers are:
 - `int`: represents arbitrarily large integers (in Python 2.x this is equivalent to C `long`);
 - `float`: floating-point numbers (equivalent to 64-bit C `double`);
 - `complex`: complex numbers.

In [15]:

```
# An integer number  
x = 123456789  
# A very large integer obtained from the one before by rising it to the 8-th power  
x ** 8
```

Out[15]: 53965948844821664748141453212125737955899777414752273389058576481

A Quick Note on Extremely Large Integers

- On Python 2.x, `sys.maxint` gives you the (maximum) integer value which your computer can work **natively** with.
- Up to `sys.maxint` your machine is able to perform arithmetic operation (e.g., addition, multiplication) in a **single** CPU instruction.
- This value corresponds to the number that can be represented using 64 bits (if your platform word's size is 64 bits, otherwise 32 bits, etc.).

A Quick Note on Extremely Large Integers: Beyond `sys.maxint`

- Just because that is what can be done in a single CPU instruction does not mean you cannot go beyond that limit!
- Python introduces **extended-precision integers** to overcome such a limitation.
- Those are "software structures" that can handle integers of any size transparently to the user by chaining them together, only limited by the memory available.
- Python 2.x keeps native integers "separate" from extended-precision ones, whilst Python 3.x treats every integer as extended-precision.

In [16]:

```
# A float number  
x = 3.645  
# A float number defined using scientific notation  
x_exp = 8.21e-4
```

A Quick Note of Floating-Point Arithmetic

- As opposed to integers, floating-point numbers have a finite-precision representation in computer hardware as base 2 (binary) fractions.
- For example, consider the decimal fraction 0.125 and the binary fraction 0.001
- Both represent the same number:
$$1 * 10^{-1} + 2 * 10^{-2} + 5 * 10^{-3} = 0 * 2^{-1} + 0 * 2^{-2} + 1 * 2^{-3}$$
- Unfortunately, most decimal fractions cannot be represented *exactly* as binary fractions.
- As such, decimal floating-point numbers are thus approximated by the binary floating-point numbers actually stored in the machine.

A Quick Note of Floating-Point Arithmetic: Issues

- No matter how many decimal digits you use, you will not get the exact representation of the fraction $1/3 = 0.333\dots$
- In the same way, no matter how many binary digits you use, the decimal value $1/10 = 0.1$ cannot be represented exactly as a binary fraction.
- In base 2, the decimal value $1/10 = 0.1$ is the infinitely repeating fraction: $0.00011001100110011\dots$
- Stop at any finite number of bits, and you get an approximation!

A Quick Note of Floating-Point Arithmetic: Example

- Suppose we want to transform a decimal number $n = 4.47$ into its corresponding yet approximated binary fraction using $k = 6$ bits of precision.
- **Step 1:** Conversion of the integer part of n (i.e., 4) to binary:
 1. $4/2$: Remainder = 0 : Quotient = 2
 2. $2/2$: Remainder = 0 : Quotient = 1
 3. $1/2$: Remainder = 1 : Quotient = 0

So, equivalent binary of integral part of decimal is **100**

A Quick Note of Floating-Point Arithmetic: Example

- **Step 2:** Conversion of the fractional part of n (i.e., .47) to binary:

1. $0.47 * 2 = 0.94$, Integral part: 0

2. $0.94 * 2 = 1.88$, Integral part: 1

3. $0.88 * 2 = 1.76$, Integral part: 1

4. $0.76 * 2 = 1.32$, Integral part: 1

5. $0.32 * 2 = 0.64$, Integral part: 0

6. $0.64 * 2 = 1.28$, Integral part: 1

So, equivalent binary of fractional part of decimal is **.011101**

A Quick Note of Floating-Point Arithmetic: Example

- **Step 3:** Combining the result of Step 1 and 2 to get the k -bit ($k = 6$) approximated binary fraction corresponding to the decimal number $n = 4.47$.

$$(4.47)_{10} = 100 + 0.011101 = (100.011101)_2$$

Operations (except for `complex`)

Operation	Result
<code>x + y</code>	sum of <code>x</code> and <code>y</code>
<code>x - y</code>	difference of <code>x</code> and <code>y</code>
<code>x * y</code>	product of <code>x</code> and <code>y</code>
<code>x / y</code>	quotient of <code>x</code> and <code>y</code>
<code>x // y</code>	floored quotient of <code>x</code> and <code>y</code>
<code>x % y</code>	remainder of <code>x / y</code>
<code>-x</code>	<code>x</code> negated
<code>+x</code>	<code>x</code> unchanged
<code>abs(x)</code>	absolute value or magnitude of <code>x</code>
<code>int(x)</code>	<code>x</code> converted to integer
<code>float(x)</code>	<code>x</code> converted to floating point
<code>complex(re, im)</code>	a complex number with real part <i>re</i> , imaginary part <i>im</i> . <i>im</i> defaults to zero.
<code>c.conjugate()</code>	conjugate of the complex number <code>c</code>
<code>divmod(x, y)</code>	the pair <code>(x // y, x % y)</code>
<code>pow(x, y)</code>	<code>x</code> to the power <code>y</code>
<code>x ** y</code>	<code>x</code> to the power <code>y</code>

Division (/): Python 2.x vs. Python 3.x

- In Python 2.x, dividing two integers always results in an `int` (C-style).
- In Python 3.x, dividing two integers always returns a `float`.
- This is fine when the result of your integer division is an integer, but it leads to quite different results when the answer is a real number!

```
# Python 2.x
# Division operator (/) always returns an int
print 4/2
2
print 3/2
1
```

```
# Python 3.x
# Division operator (/) always returns a float
print(4/2)
2.0
print(3/2)
1.5
```

Integer Division in Python 3.x

- To get C-style integer division in Python 3.x, use the floor division operator `//`:

```
print(3//2)  
1
```

Sequences

Sequence Types

- Built-in sequences can be either *immutable* or *mutable*.
- *Immutable* sequence types are:
 - `str`
 - `bytes`
 - `tuple`
- *Mutable* sequence types are:
 - `byte array`
 - `list`

Operations

The operations in the following table are supported by most sequence types, both *mutable* and *immutable*.

Operation	Result
<code>x in s</code>	True if an item of <i>s</i> is equal to <i>x</i> , else False
<code>x not in s</code>	False if an item of <i>s</i> is equal to <i>x</i> , else True
<code>s + t</code>	the concatenation of <i>s</i> and <i>t</i>
<code>s * n</code> or <code>n * s</code>	equivalent to adding <i>s</i> to itself <i>n</i> times
<code>s[i]</code>	<i>i</i> th item of <i>s</i> , origin 0
<code>s[i:j]</code>	slice of <i>s</i> from <i>i</i> to <i>j</i>
<code>s[i:j:k]</code>	slice of <i>s</i> from <i>i</i> to <i>j</i> with step <i>k</i>
<code>len(s)</code>	length of <i>s</i>
<code>min(s)</code>	smallest item of <i>s</i>
<code>max(s)</code>	largest item of <i>s</i>
<code>s.index(x[, i[, j]])</code>	index of the first occurrence of <i>x</i> in <i>s</i> (at or after index <i>i</i> and before index <i>j</i>)
<code>s.count(x)</code>	total number of occurrences of <i>x</i> in <i>s</i>

Strings: Type `str` (*immutable*)

String Definition

- You can write *string literals* using either single quotes `'` or double quotes `"`.
- Similarly, multiline strings with line breaks must be enclosed by triple quotes, either `' ' '` or `" " "`.

In [17]:

```
s = 'This is a single-quoted string'
t = "This is a double-quoted string"
u = 'This is a single-quoted string with "double quotes" inside'
v = "This is a double-quoted string with 'single quotes' inside"
w = 'This is a single-quoted string with \'escaped single quotes\' inside'
x = "This is a double-quoted string with \"escaped double quotes\" inside"
```

In [18]:

```
1 m_s = '''  
2 This is  
3 a multiline string  
4 enclosed by triple single quotes  
5 '''  
6 m_t = """  
7 This is  
8 a multiline string  
9 enclosed by triple double quotes  
10 """
```

```
In [19]: 1 # Count how many lines the string above is made of  
2 # You might expect the result being 3, instead the '\n' character  
3 # right after the opening and closing triple quotes counts as well  
4len(m_s.split('\n'))
```

Out[19]: 5

Properties

- Python 3.x strings (`str`) are **immutable** sequences of Unicode **code points**.
- [Unicode](#) is a standard mapping between each character of every language to a unique number (**code point**) [to support non-ASCII characters].
- Unicode defines 1,114,112 code points, which are denoted by (hexadecimal) numbers in the range of `U+000000 - U+10FFFF`.
- In Python 2.x, `str` instead refers to a sequence of **bytes** and there is a dedicated type `unicode` for representing Unicode code points.
- You **cannot** modify a string without creating a new one.

In [20]:

```
s = 'This is a string'
# Try to access the 7-th character of the sequence (index is 0-based)
# and change it to a different character
s[6] = 'z'
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-20-052ef33b03de> in <module>()
      2 # Try to access the 7-th character of the sequence (index is 0-based)
      3 # and change it to a different character
----> 4 s[6] = 'z'
```

```
TypeError: 'str' object does not support item assignment
```

In [21]:

```
# This will actually create a new, modified string object  
new_s = s.replace('string', 'new string')  
print(new_s)
```

This is a new string

String Concatenation

- It is often very useful to be able to combine strings into a new string.
- This can be done with the plus sign (+), which is the *operator* used to concatenate two (or more) strings into one.
- You can use as many plus signs as you want in composing messages.

In [22]:

```
a = 'This is the first string.'  
b = 'This is the second string.'  
c = 'This is the third string.'  
# Concatenating them all and interleave each string with a blank character  
print(a + ' ' + b + ' ' + c)
```

This is the first string. This is the second string. This is the third string.

Quiz Time

Concatenating more than 2 strings using the '+' operator doesn't scale well and might be highly inefficient when the number of strings to concatenate becomes larger. **Why?**

Answer

Because for each concatenation (i.e., for each pair of strings to concatenate) a *new* string object is created (allocated) and all the previous strings have to be first copied into the newly allocated space for result

Suppose you have n strings (therefore $n - 1$ concatenations), each string of length l : you'll copy $2l$ characters for the first concatenation (i.e., l from the first and l from the second string), plus $3l$ the second concatenation, plus $4l$ the third concatenation, and so on and so forth.

Overall:

$$l * \sum_{i=2}^n i = l * \left[\frac{n(n+1)}{2} - 1 \right],$$

which is, indeed, $O(n^2)$.

[As of Python 2.4, the CPython implementation avoids creating a new string object when using $a += b$ or $a = a + b$, but this optimization is both fragile and not portable.]

More Efficient String Concatenation

- Use `" ".join([a, b, c])`

In [23]:

```
print(" ".join([a, b, c]))  
# alternatively, use a different separator from whitespace (e.g., '\n')  
print("\n".join([a, b, c]))
```

This is the first string. This is the second string. This is the third string.

String Formatting

- String templating or formatting is another important topic.
- The number of ways to do so has expanded with the advent of Python 3
- String objects have a `format` method which can be used to substitute formatted arguments into the string, producing a new string.
- More information can be found on Python official [documentation](#).

In [24]:

```
1 # Suppose you have multiple strings that are made of some fixed portion
2 # as well as some variable portions that all adhere to a specific formatting pattern.
3 # Let's define the following formatting pattern
4 template = '{0:.2f} {1:s} are worth US${2:d}'
5
6 # In the above template string:
7 # {0:.2f} means to format the first argument as a floating point number with 2 decimals.
8 # {1:s} means to format the 2nd argument as a string.
9 # {2:d} means to format the 3rd argument as an exact integer.
10
11 # We perform parameter substitution on the template defined above using the format method
12 print(template.format(4.5560, 'Argentine Pesos', 1))
```

4.56 Argentine Pesos are worth US\$1

In [25]:

```
# If the order of the arguments of .format is the same of that expected by template  
# you can omit the indices: 0, 1, 2, etc.  
template = '{:.2f} {:s} are worth US${:d}'  
4  
# We perform parameter substitution on the template defined above using the format method  
print(template.format(4.5560, 'Argentine Pesos', 1))
```

4.56 Argentine Pesos are worth US\$1

In [26]:

```
1 # Otherwise, you could specify a different order in the template w.r.t. the one of .format
2 # BE CAREFUL WITH THIS APPROACH!
3 template = '{2:.2f} {0:s} are worth US${1:d}'
4
5 # We perform parameter substitution on the template defined above using the format method
6 print(template.format(4.5560, 'Argentine Pesos', 1))
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-26-150fe06fb6c3> in <module>()
      4
      5 # We perform parameter substitution on the template defined above using the format
      method
----> 6 print(template.format(4.5560, 'Argentine Pesos', 1))

ValueError: Unknown format code 's' for object of type 'float'
```

In [27]:

```
# Otherwise, you could specify a different order in the template w.r.t. the one of .format  
# BE CAREFUL WITH THIS APPROACH!  
template = '{2:.2f} {0:s} are worth US${1:d} '  
4  
# We perform parameter substitution on the template defined above using the format method  
print(template.format('Argentine Pesos', 1, 4.5560))
```

4.56 Argentine Pesos are worth US\$1

Unicode vs. Byte Strings

Python 2.x

- In Python 2.x there are 2 distinct types of strings:
 - `str` --> refers to sequence of bytes;
 - `unicode` --> refers to sequence of Unicode code points.
- Depending on the **character encoding** used (e.g., UTF-8, ISO 8859-1, etc.) the same code point is possibly mapped to a different sequence of bytes.

Unicode as an Abstraction of Text

letter	Unicode Code Point
ć	\u0107

Byte Encodings			
letter	UTF-8	UTF-16	Shift-JIS
ć	\xc4\x87	\x07\x01	\x85\xc9

From Byte to Unicode String in Python 2.x

- To convert a Python 2.x byte string object (`str`) into its corresponding Unicode object (`unicode`) you need to call the `decode(character_encoding)` method (assuming you know `character_encoding`, e.g., UTF-8)

```
# Assuming this is a UTF-8 encoded Python 2.x str  
s = 'This is a UTF-8 byte string' # s has type str  
u_s = s.decode("UTF-8") # u_s has type unicode
```

From Unicode to Byte String in Python 2.x

- Every time you have to serialize out your string you need to transform it into a sequence of bytes!
- To do so, use the `encode(character_encoding)` method.
- **Warning:** Not every Unicode sequence can be encoded by every character encoding! For example, ASCII character encoding can only encode Unicode sequences representing ASCII characters.
- Here is a comprehensive [reference](#) to all we have been discussing so far.

Luckily, We Use Python 3.x!

- Since Python 3.0, Unicode has become the first-class string type to enable more consistent handling of ASCII and non-ASCII text.
- Now the type `str` refers to Unicode **not** to bytes!
- There is however a specific type `bytes` to explicitly indicate sequence of bytes.

In [28]:

```
1 print('***** From Unicode string to byte string *****')
2 # This is a Unicode string containing non-ASCII character
3 s = 'Barça'
4 # This statement prints the type associated with s
5 print(type(s))
6 # We still can convert this Unicode string
7 # to its UTF-8 bytes representation using the encode method:
8 s_utf8 = s.encode("utf-8")
9 print(s_utf8)
10 print(type(s_utf8))
11 # If we try to encode our Unicode sequence to ASCII encoding...
12 s_ascii = s.encode("ascii")
```

```
***** From Unicode string to byte string *****
```

```
<class 'str'>
```

```
b'Bar\xc3\xa7a'
```

```
<class 'bytes'>
```

```
-----
UnicodeEncodeError                                Traceback (most recent call last)
<ipython-input-28-b1047bf72ce8> in <module>()
      10 print(type(s_utf8))
      11 # If we try to encode our Unicode sequence to ASCII encoding...
--> 12 s_ascii = s.encode("ascii")
```

```
UnicodeEncodeError: 'ascii' codec can't encode character '\xe7' in position 3: ordinal not
in range(128)
```

In [29]:

```
1 print('***** From byte string to Unicode string *****')
2 # Assuming you know the Unicode encoding of a bytes object,
3 # you can still go back using the decode method:
4 s_unicode = s_utf8.decode("utf-8")
5 print(s_unicode)
6 print(type(s_unicode))
7 # Again, if we try to decode the byte sequence with a different encoding
8 # than the one actually used to serialize the Unicode sequence...
9 s_unicode = s_utf8.decode("ascii")
```

```
***** From byte string to Unicode string *****
```

```
Barça
```

```
<class 'str'>
```

```
-----
UnicodeDecodeError                                Traceback (most recent call last)
<ipython-input-29-adbff3796966> in <module>()
      7 # Again, if we try to decode the byte sequence with a different encoding
      8 # than the one actually used to serialize the Unicode sequence...
----> 9 s_unicode = s_utf8.decode("ascii")
```

```
UnicodeDecodeError: 'ascii' codec can't decode byte 0xc3 in position 3: ordinal not in range(128)
```

Not Everything Needs To Be UTF-8-encoded!

- While it is become preferred to use UTF-8 for any encoding, for historical reasons you may encounter data in any number of different encodings:
 - UTF-16
 - ISO 8859-1 (latin1)
 - Windows-1252 (CP-1252)
 - ...

In [30]:

```
print(s.encode("utf-16"))  
print(s.encode("iso-8859-1"))  
print(s.encode("windows-1252"))
```

b'\xff\xfeB\x00a\x00r\x00\xe7\x00a\x00'

b'Bar\xe7a'

b'Bar\xe7a'

Bytes: Type `bytes` (*immutable*)

Sometimes You Just Need Bytes!

- Especially while working with binary files (i.e., files containing sequence of bytes).
- A sequence of bytes is a sequence of integers in the range of **0–255** (only available in Python 3.x).
- You may not want to **decode** those sequence of bytes to Unicode sequence of chars!
- Note however that you can define your own byte literals by prefixing a string with **b**:

```
byte_string = b'This is a byte string'
```

ByteArray: Type `bytearray` (*mutable*)

Properties

- This built-in data type corresponds to *mutable* **bytes**.
- It is only available in Python 3.x.

Summary

- Built-in data types:
 - `bool` and `NoneType` (`None`)
 - numeric: `int`, `float`, `complex` (*immutable*)
 - sequences: `str`, `bytes` (*immutable*), `bytearray` (*mutable*)
 - More built-in data types in the next lecture!