

Fundamentals of Information Systems

Python Programming (for Data Science)

Master's Degree in Data Science

Gabriele Tolomei

gtolomei@math.unipd.it

University of Padua, Italy

2018/2019

December, 3 2018

Lecture 11: A Machine Learning Primer

What is Machine Learning?

- There exist several definitions of *what* is **Machine Learning (ML)**:
 1. "A computer program is said to **learn from experience** E with respect to some class of tasks T and performance measure P if its performance at tasks in T , as measured by P , improves with experience E ." ([Tom Mitchell](#))
 2. "ML is an application of artificial intelligence (AI) that provides systems the ability to **automatically learn** and improve from experience (i.e., *observed data*) **without being explicitly programmed**" ([source](#)).
 - ...

Just Take a Step Back

- Traditionally, we design computer programs to solve specific tasks which would otherwise require a huge human effort.
- For example, consider the task of creating a vocabulary of words from a collection of text documents.
- To do so, we need to be able to build an **algorithm** which solves the task we aim to perform.
- In other words, we have to *mathematically encode* our task into a **computable function** which takes some **input** (e.g., a collection of text documents) and returns some **output** (e.g., a vocabulary of words).

Different Types of Tasks

- Some tasks, however, simply **cannot** be pinpointed down mathematically, as the function they encode to is just too hard to be implemented in a traditional computer program.
- Consider how would you possibly write a computer program, which given as input an image gets as output a boolean answer that tells whether the image contains or not a bird...
- **Idea:** Access to a collection of, say, 1M images each one with a **label** associated (e.g., **bird/no_bird**) and output a computer program that detects birds in images.

xkcd Tells Us Better!



IN CS, IT CAN BE HARD TO EXPLAIN
THE DIFFERENCE BETWEEN THE EASY
AND THE VIRTUALLY IMPOSSIBLE.

([Image Source](#))

Taxonomy of ML

- Roughly speaking there are two main **categories** of ML tasks:
 1. **Supervised Learning:** Given a set of **labeled** examples, **predict** the labels of *new and unseen* examples
 2. **Unsupervised Learning:** Given a set of examples, **find structure** in the data (e.g., *clusters, subspaces, manifolds*)

Supervised Learning

- We can further divide supervised learning into two sub-categories, depending on the **label** we want to predict.
 - **Regression:** if the label takes on *continuous* values
 - **Classification:** if the label takes on *discrete* values

Supervised Learning: Regression vs. Classification

- Regression:
 - value of a house
 - probability of click on a link to a web page
- Classification:
 - spam vs. non-spam emails (**binary classification**)
 - movie ratings on a 1 to 5 scale (**k-ary or multiclass classification**)

The General (Supervised) ML Pipeline

- 0) Be sure your problem needs *actually* to be tackled using ML!
- 1) **Data Collection:** Get **labeled data** from the domain of your interest (may be the hardest part!)
- 2) **Feature Engineering:** Represent your data with a "machine-friendly" format
- 3) **Model Training:** Build one (or more) **learning models** (more on this later)
- 4) **Model Selection/Evaluation:** Pick the best-performing model according to some **quality metrics**

1. Data Collection

Motivation

- Any ML technique requires data to operate on!
- Supervised Learning, in particular, needs **labeled data**, which may be even harder to get (e.g., a set of emails with **spam/non_spam** tags).
- Most of the time, this step involves combining multiple and possibly heterogeneous data sources.

2. Feature Engineering

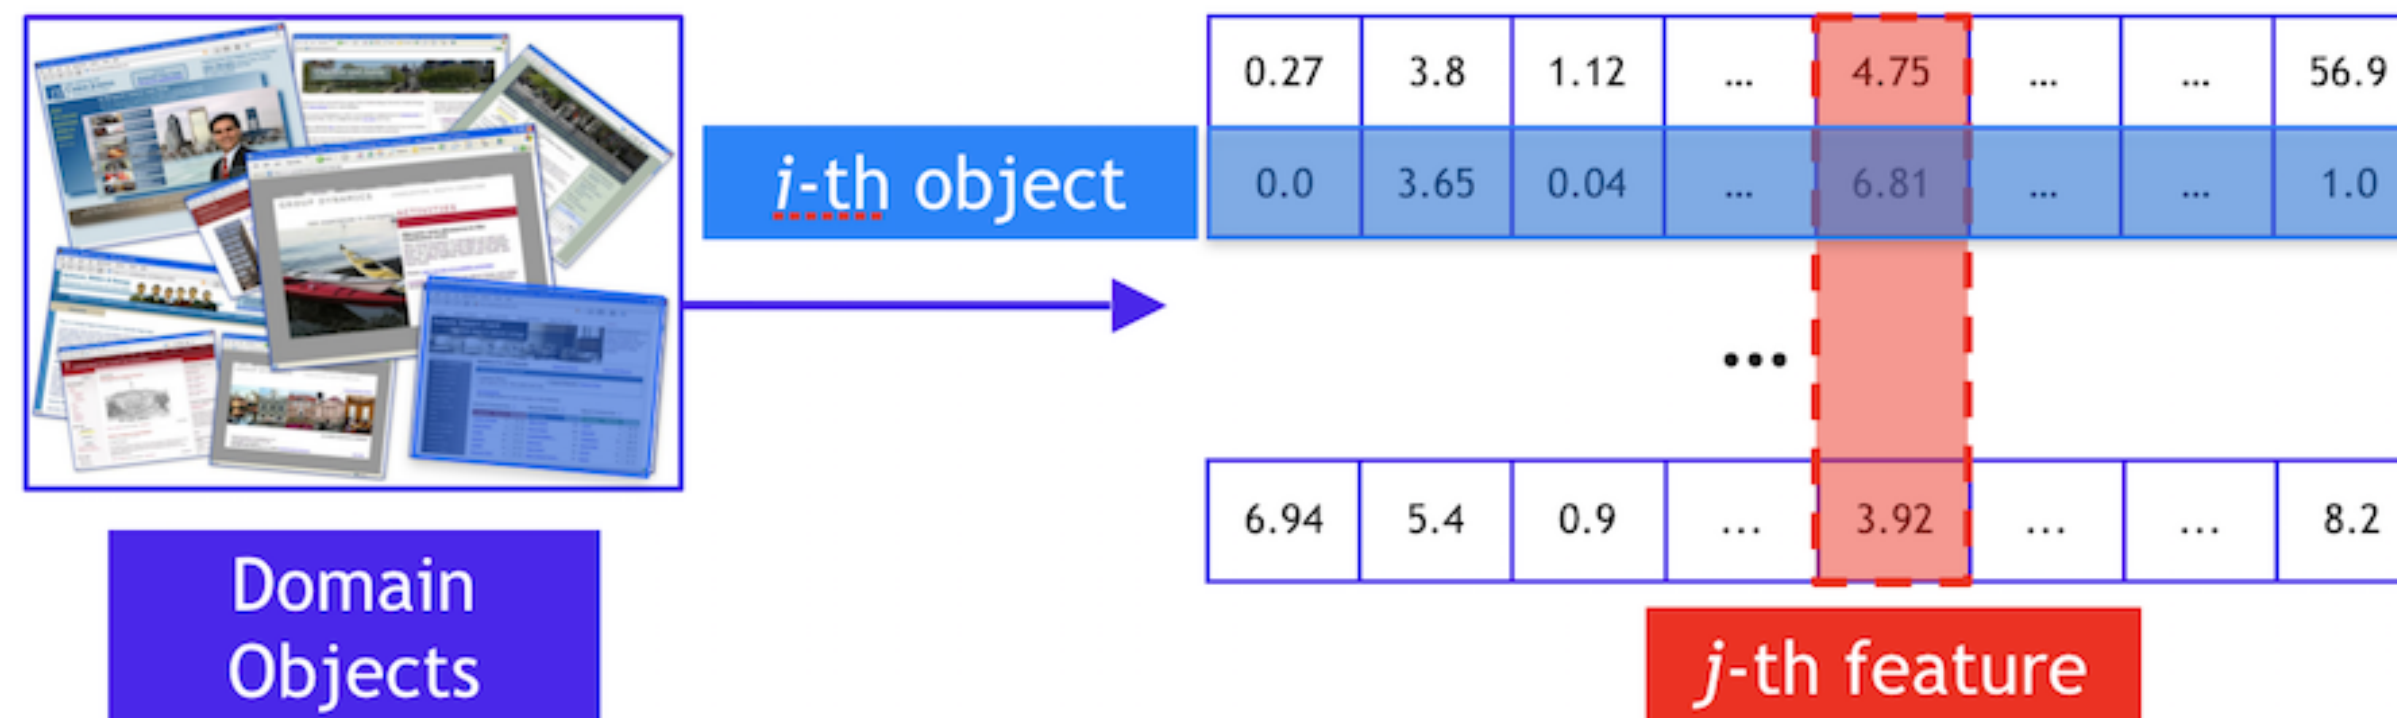
Motivation

- Collected data must be encoded to a **machine-readable** format.
- Each object of our problem domain is transformed into an n -dimensional **feature vector**.
- Each **feature** captures a specific property of the object we believe it is useful to "separate" different objects.
- In the following, we refer to **instance** (or **example**) to identify the feature vector of an object **and** its associated label.

Characteristics of Features

- Feature values can be either **categorical** (e.g., *gender*) or **continuous** (e.g., *weight*).
- A feature can be derived *locally* from a single object (e.g., the *annual salary* of a person)
- Or it can result from more complex computation on the whole data collection (e.g., the *tf-idf* score of a word in a document computed against a corpus of documents).

Example: *spam* vs. *non-spam* emails



Challenges and Solutions

- **Problem:** Collected data is far from being perfect!
- Many challenges need to be addressed before moving down to next stages of the ML pipeline.
- Overall, this is accomplished with an intermediate **Data Preprocessing** step.

Handling Missing Values

Challenge	Description	Solution
Missing values	A feature value may not be available for one or more instances	Replace missing values with the median (<i>continuous</i>) or the mode (<i>categorical</i>) of the existing values

Managing Data Sparsity

Challenge	Description	Solution
Sparsity	Most of the instances contain just a small subset of the features	Use “sparse-friendly” data structures (e.g., DOK)

Handling Outliers

Challenge	Description	Solution
Outliers	One or more instances have out-of-range values for one or more features	Retention vs. Exclusion (<i>trimming</i> or <i>Winsorising</i>)

Mixing Categorical and Continuous Feature Values

Challenge	Description	Solution
Mix of continuous and discrete values	Feature set contains both numerical and categorical values	Transform categorical features using <i>one-hot encoding</i>

Different Feature Scales

Challenge	Description	Solution
Multiple feature magnitudes	Feature set contains very wide range of values	Standardization (min-max, z-scores)

Managing Imbalance

Challenge	Description	Solution
Class imbalance	Instances labeled with the class of interest represents a tiny fraction of the total	Over-/Under-sampling, cost-sensitive learning


Relationship between Features

Challenge	Description	Solution
Strong multicollinearity	Linear relationship between one or more features	Dimensionality reduction (PCA)

A Bit of Notation

- We denote by $\mathcal{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m)\}$ our resulting labeled dataset.
- This is made of m **instances**; the i -th instance is represented by the tuple (\mathbf{x}_i, y_i) , where:
 - $\mathbf{x}_i \in \mathbb{R}^n$ is an n -dimensional **feature vector**, i.e.,
$$\mathbf{x}_i = (x_{i,1}, \dots, x_{i,n})^T$$
 - $y_i \in \mathbb{R}$ if the label is **continuous (regression)** or
 $y_i \in \{1, 2, \dots, k\}$ if the label is **categorical** and takes on k values (**k-ary classification**)


Labeled Dataset



A Bit of Notation

- We indicate by $X \in \mathbb{R}_{m,n}$ the m -by- n **feature matrix** resulting from the m n -dimensional feature vectors.
- Similarly, we denote by Y the m -dimensional vector containing instance labels.

Labeled Dataset



3. Model Training

Intuition

- **Idea:** There exists an **unknown target function** f , which maps element of X to elements of Y .

$$f = X \mapsto Y$$

- Unfortunately, we cannot just write down an algorithm that implements f .
- Still, we can try to **learn** f , namely to find another function h^* which approximates the *true* f ($h^* \approx f$) using the data \mathcal{D} we observed.
- We refer to h^* as our **hypothesis** (used to approximate f).

How Do We Select h^* From?

- We choose h^* from a family of functions \mathcal{H} , called the **hypothesis space**.
- To actually pick h^* , we need to further specifying two components:
 - the **loss function** (a.k.a. **cost function**) denoted by ℓ
 - the **learning algorithm** denoted by \mathcal{A}

The Loss Function ℓ

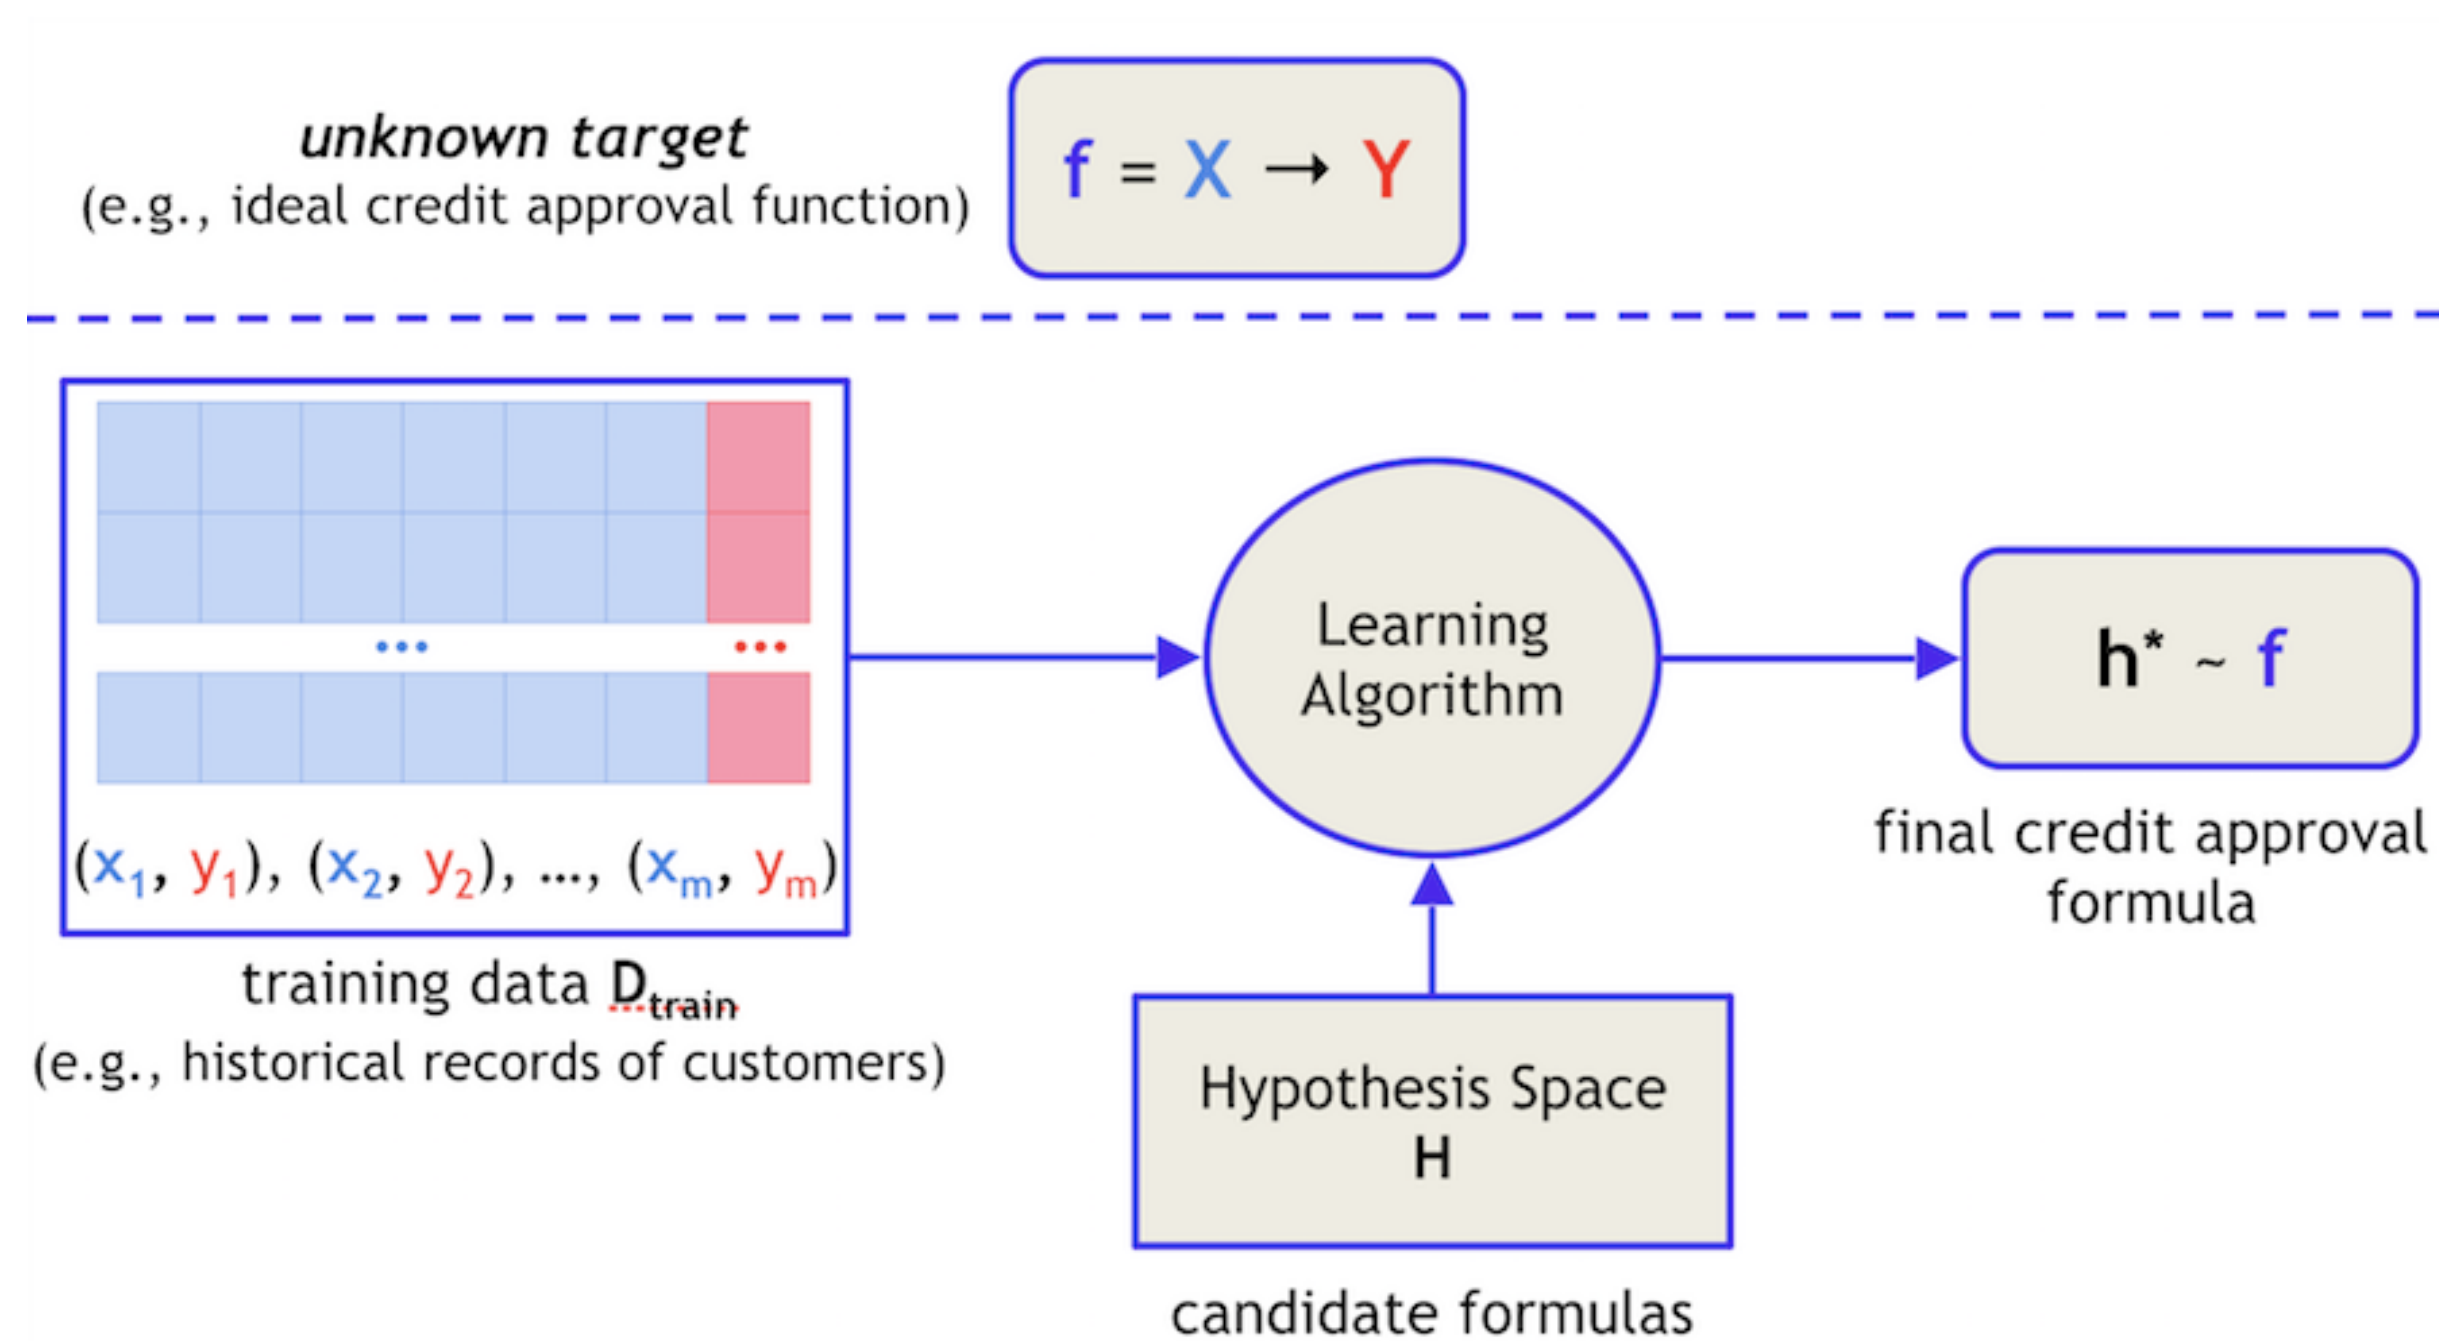
- Measures the error we would make if a hypothesis h is used instead of the true f .
- Such an error can be computed *only* on the data we observed (i.e., on \mathcal{D}).
- Therefore, it depends on the hypothesis *and* the dataset, i.e.,
 $\ell : \mathcal{H} \times \mathcal{D} \mapsto \mathbb{R}$.
- We should use this **in-sample error** (a.k.a. **empirical loss**) as an *estimate* of the **out-of-sample error** (a.k.a. **expected loss** or **risk**).

The Learning Algorithm \mathcal{A}

- Defines the strategy we use to search the hypothesis space \mathcal{H} for picking our **best** hypothesis $h^* \in \mathcal{H}$.
- Here, "**best**" means the hypothesis that **minimizes** the loss function on the observed data (**Empirical Risk Minimization**).
- In other words, among all the hypotheses specified by \mathcal{H} , the learning algorithm will pick the one that minimizes ℓ .

$$h^* = \operatorname{argmin}_{h \in \mathcal{H}} \ell(h, \mathcal{D})$$

One-Slide ML



The Hypothesis Space \mathcal{H}

- Intuitively, the larger the hypothesis space:
 - the larger will be the set of functions that can be represented (+)
 - the harder will be for the learning algorithm to pick the best h^* (-)
- **Trade-off:** Put some constraints on \mathcal{H} , e.g., limit the search space \mathcal{A} has to explore only to **linear functions**.

Example of \mathcal{H} : Linear Models

- The underlying assumption here is that there exists a **linear relationship** between X (**features**) and Y (**class label/value**).

$$\mathcal{H} = \{h_{\theta} : X \mapsto Y \mid h_{\theta}(\mathbf{x}) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n\}$$

- θ is called the **vector of parameters** (of the linear model).
- $x_0 = 1$ by convention.
- Among all the possible instantiations of θ , the learning algorithm will pick θ^* as the one which minimizes the **loss function** ℓ computed on \mathcal{D} .

A Possible Loss Function: *Mean Squared Error* (MSE)

- **MSE** is just a possible (mathematically convenient) choice of loss function ℓ .
- **MSE** measures the average error when the true target f is replaced with a hypothesis $h_{\theta} \in \mathcal{H}$ on the observed data \mathcal{D} .

$$\begin{aligned}\ell(h_{\theta}, \mathcal{D}) &= \text{MSE}(h_{\theta}, \mathcal{D}) \\ &= \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(\mathbf{x}_i) - f(\mathbf{x}_i) \right)^2 \\ &= \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(\mathbf{x}_i) - y_i \right)^2\end{aligned}$$

- Each term of the above summation, i.e., $\left(h_{\theta}(\mathbf{x}_i) - y_i \right)^2$, represents the **squared error** w.r.t. a single instance in \mathcal{D} .

The Learning Algorithm (Revisited)

- Reduces the learning problem to an **optimization problem**.
- Among all the possible $h_{\theta} \in \mathcal{H}$, it picks $h^* = h_{\theta^*}$, so as to minimize the loss function.

$$h^* = \operatorname{argmin}_{\theta} \{ \ell(h_{\theta}, \mathcal{D}) \}$$

- In particular, if $\ell = \text{MSE}$:

$$h^* = \operatorname{argmin}_{\theta} \{ \text{MSE}(h_{\theta}, \mathcal{D}) \} = \operatorname{argmin}_{\theta} \left\{ \frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(\mathbf{x}_i) - y_i \right)^2 \right\}$$

Minimum and Maximum of a Function: First Derivative

- Whenever we want to find the minimum (maximum) of a real-valued function $f : \mathbb{R} \mapsto \mathbb{R}$, we look for **stationary points**, i.e., the set of points $S = \{(x, f(x)) \mid f'(x) = 0\}$ where the first derivative is 0.
- Note that depending on the shape of the function f , a stationary point doesn't necessarily correspond to a local minimum (maximum) but it may be a **point of inflection**.
- To investigate better what happens at a specific stationary point, we need to check the second derivative $f''(x)$.

Minimum and Maximum of a Function: Second Derivative

- Given a stationary point $(x_s, f(x_s)) \in S$:
 - if $f''(x_s) > 0$ then $(x_s, f(x_s))$ is a **local minimum**
 - if $f''(x_s) < 0$ then $(x_s, f(x_s))$ is a **local maximum**
 - if $f''(x_s) = 0$ then $(x_s, f(x_s))$ may be an **inflection point**

Minimum and Maximum of a Convex/Concave Function

- Any local minimum (maximum) of a **convex** (**concave**) function is also a **global** minimum (maximum).
- If we know the function is **convex** (**concave**) finding the minimum (maximum) can be done just by computing the first derivative and set it to 0.
- In the case of a multivariate function, this generalizes to compute the **gradient** (∇) of the function and set it to 0.

The Gradient ∇

- The gradient of a multivariate function $f : \mathbb{R}^n \mapsto \mathbb{R}$ is the n -dimensional vector of the **partial derivatives** of the function w.r.t. each of its variable.

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right)$$


- Solving $\nabla f = \mathbf{0}$ means finding an n -dimensional vector resulting from the following:

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right) = \underbrace{(0, 0, \dots, 0)}_n = \mathbf{0}$$

MSE(h_{θ} , \mathcal{D}) is a Convex Function

- Each term of the summation of MSE(h_{θ} , \mathcal{D}) is a multivariate linear function of the model parameters $\theta = (\theta_0, \theta_1, \dots, \theta_n)$.
- Linear functions are **convex**, and so does the square of a convex function; also, the sum of convex functions is again convex.
- Convex functions have a **unique local minimum**, which therefore happens to be the **global minimum**.

Visualize The Optimization Problem



Computing the Gradient of MSE

- The following is the expression of the gradient of $\text{MSE}(h_{\theta}, \mathcal{D})$:

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \nabla \left[\frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(\mathbf{x}_i) - y_i \right)^2 \right]$$

Computing the Gradient of MSE

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \nabla \left[\frac{1}{m} \sum_{i=1}^m \left(h_{\theta}(\mathbf{x}_i) - y_i \right)^2 \right]$$

- Due to the **scalar multiple rule** and **sum rule** of the derivative, i.e., $\frac{\partial f}{\partial t}(\alpha t) = \alpha \frac{\partial f}{\partial t}$ ($\alpha \in \mathbb{R}$, constant) and $\frac{\partial f}{\partial t} \left(\sum t \right) = \sum \left(\frac{\partial f}{\partial t} \right)$, we can rewrite the above as follows:

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \frac{1}{m} \left[\sum_{i=1}^m \nabla \left(h_{\theta}(\mathbf{x}_i) - y_i \right)^2 \right]$$

Computing the Gradient of MSE (single instance)

- To make things easier, let's consider for a moment the special case where \mathcal{D} contains just a **single** instance, i.e., $\mathcal{D} = \{(\mathbf{x}, y)\}$. Therefore:

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \nabla \left(h_{\theta}(\mathbf{x}) - y \right)^2$$

- Due to the **power rule** and **chain rule** of the derivative, i.e., $\frac{\partial f(t^{\alpha})}{\partial t} = \alpha t^{\alpha-1} \frac{\partial f(g(t))}{\partial t} = \frac{\partial f(g(t))}{\partial g(t)} \cdot \frac{\partial g(t)}{\partial t}$, we can rewrite the above as follows:

$$= 2 \left(h_{\theta}(\mathbf{x}) - y \right) \nabla \left(h_{\theta}(\mathbf{x}) - y \right)$$

Computing the Gradient of MSE (single instance)

$$\begin{aligned}\nabla \left(h_{\boldsymbol{\theta}}(\mathbf{x}) - y \right) &= \nabla \left(\theta_0 + \theta_1 x_1 + \dots + \theta_n x_n - y \right) \\ &= \left(\frac{\partial(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n - y)}{\partial \theta_0}, \frac{\partial(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n - y)}{\partial \theta_1}, \dots \right. \\ &\quad \left. \frac{\partial(\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n - y)}{\partial \theta_n} \right) = \left(x_0, x_1, \dots, x_n \right) = \mathbf{x}\end{aligned}$$

- Overall:

$$\nabla \text{MSE}(h_{\boldsymbol{\theta}}, \mathcal{D}) = 2 \underbrace{\left(h_{\boldsymbol{\theta}}(\mathbf{x}) - y \right)}_{\text{scalar}} \cdot \underbrace{\mathbf{x}}_{(n+1)\text{-dimensional vector}}$$

Computing the Gradient of MSE (single instance)

- We can rewrite the formula above using **vectorized notation**, noticing that $h_{\theta}(\mathbf{x}) = \theta_0 x_0 + \theta_1 x_1 + \dots \theta_n x_n = \boldsymbol{\theta}^T \mathbf{x}$.
- $\boldsymbol{\theta}^T \mathbf{x}$ is obviously a scalar which results from the dot product of a 1-by- $(n + 1)$ row vector $\boldsymbol{\theta}^T$ and an $(n + 1)$ -by-1 column vector \mathbf{x} .

$$\nabla \text{MSE}(h_{\theta}, D) = 2(\boldsymbol{\theta}^T \mathbf{x} - y) \cdot \mathbf{x}$$

Computing the Gradient of MSE (single instance)

- The product of a vector by a scalar is applied element-wise, therefore:

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \begin{bmatrix} 2(\boldsymbol{\theta}^T \mathbf{x} - y) \cdot x_0 \\ 2(\boldsymbol{\theta}^T \mathbf{x} - y) \cdot x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \mathbf{x} - y) \cdot x_n \end{bmatrix} = \begin{bmatrix} 2(\boldsymbol{\theta}^T \mathbf{x} - y) \\ 2(\boldsymbol{\theta}^T \mathbf{x} - y) \cdot x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \mathbf{x} - y) \cdot x_n \end{bmatrix}$$

as we set $x_0 = 1$.

- Note that $\nabla \text{MSE}(h_{\theta}, \mathcal{D})$ is an $(n + 1)$ -dimensional vector, as expected.

Setting the Gradient of MSE to 0 (single instance)

- To find the local minimum (which is also global) of $\text{MSE}(h_{\theta}, \mathcal{D})$, we set:

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \begin{bmatrix} 2(\boldsymbol{\theta}^T \mathbf{x} - y) \\ 2(\boldsymbol{\theta}^T \mathbf{x} - y) \cdot x_1 \\ \vdots \\ 2(\boldsymbol{\theta}^T \mathbf{x} - y) \cdot x_n \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \mathbf{0}$$

- Therefore, we have to solve a system of $n + 1$ linear equations of $n + 1$ variables, each equation of the form $2(\boldsymbol{\theta}^T \mathbf{x} - y) \cdot x_j = 0$,
 $j \in \{0, 1, \dots, n\}$

Computing the Gradient of MSE (m instances)

- Going back to the overall **MSE** in the general case where \mathcal{D} contains m instances:

$$\begin{aligned}\nabla \text{MSE}(h_{\theta}, \mathcal{D}) &= \frac{1}{m} \left\{ \left[2 \left(h_{\theta}(\mathbf{x}_1) - y_1 \right) \nabla \left(h_{\theta}(\mathbf{x}_1) - y_1 \right) \right] + \right. \\ &\quad \dots \\ &\quad \left. + \left[2 \left(h_{\theta}(\mathbf{x}_m) - y_m \right) \nabla \left(h_{\theta}(\mathbf{x}_m) - y_m \right) \right] \right\} \\ &= \frac{2}{m} \left[\sum_{i=1}^m \left(h_{\theta}(\mathbf{x}_i) - y_i \right) \nabla \left(h_{\theta}(\mathbf{x}_i) - y_i \right) \right]\end{aligned}$$

Computing the Gradient of MSE (m instances)

- Putting all together:

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \frac{2}{m} \left[\sum_{i=1}^m \underbrace{\left(h_{\theta}(\mathbf{x}_i) - y_i \right)}_{\text{scalar}} \cdot \underbrace{\mathbf{x}_i}_{(n+1)\text{-dimensional vector}} \right]$$

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \begin{bmatrix} \frac{2}{m}(\boldsymbol{\theta}^T \mathbf{x}_1 - y_1) \cdot x_{1,0} + \dots + \frac{2}{m}(\boldsymbol{\theta}^T \mathbf{x}_m - y_m) \cdot x_{m,0} \\ \frac{2}{m}(\boldsymbol{\theta}^T \mathbf{x}_1 - y_1) \cdot x_{1,1} + \dots + \frac{2}{m}(\boldsymbol{\theta}^T \mathbf{x}_m - y_m) \cdot x_{m,1} \\ \vdots \\ \frac{2}{m}(\boldsymbol{\theta}^T \mathbf{x}_1 - y_1) \cdot x_{1,n} + \dots + \frac{2}{m}(\boldsymbol{\theta}^T \mathbf{x}_m - y_m) \cdot x_{m,n} \end{bmatrix}$$

Computing the Gradient of MSE (m instances)

- Note that we have introduced a second subscript to be able to index each feature's instance within \mathcal{D} , i.e., $x_{i,j}$.
- Moreover, we have set all $x_{i,0} = 1$ ($i \in \{1, \dots, m\}$).

$$\nabla \text{MSE}(h_{\theta}, \mathcal{D}) = \frac{2}{m} \begin{bmatrix} (\theta^T \mathbf{x}_1 - y_1) + \dots + (\theta^T \mathbf{x}_m - y_m) \\ (\theta^T \mathbf{x}_1 - y_1) \cdot x_{1,1} + \dots + (\theta^T \mathbf{x}_m - y_m) \cdot x_{m,1} \\ \vdots \\ (\theta^T \mathbf{x}_1 - y_1) \cdot x_{1,n} + \dots + (\theta^T \mathbf{x}_m - y_m) \cdot x_{m,n} \end{bmatrix}$$

Setting the Gradient of MSE to 0 (m instances)

$$\underbrace{\frac{2}{m} \begin{bmatrix} (\boldsymbol{\theta}^T \mathbf{x}_1 - y_1) + \dots + (\boldsymbol{\theta}^T \mathbf{x}_m - y_m) \\ (\boldsymbol{\theta}^T \mathbf{x}_1 - y_1) \cdot x_{1,1} + \dots + (\boldsymbol{\theta}^T \mathbf{x}_m - y_m) \cdot x_{m,1} \\ \vdots \\ (\boldsymbol{\theta}^T \mathbf{x}_1 - y_1) \cdot x_{1,n} + \dots + (\boldsymbol{\theta}^T \mathbf{x}_m - y_m) \cdot x_{m,n} \end{bmatrix}}_{\nabla \text{MSE}(h_{\boldsymbol{\theta}}, D)} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \mathbf{0}$$

- Again, we have to solve a system of $n + 1$ linear equations of $n + 1$ variables, each equation of the form

$$\frac{2}{m} \left[(\boldsymbol{\theta}^T \mathbf{x}_1 - y_1) \cdot x_{1,j} + \dots + (\boldsymbol{\theta}^T \mathbf{x}_m - y_m) \cdot x_{m,j} \right] = 0,$$

$$j \in \{0, 1, \dots, n\}$$

Vectorized Form of the Gradient of MSE

$$\underbrace{\begin{bmatrix} x_{1,0}, x_{1,1}, \dots, x_{1,n} \\ x_{2,0}, x_{2,1}, \dots, x_{2,n} \\ \vdots \\ x_{m,0}, x_{m,1}, \dots, x_{m,n} \end{bmatrix}}_{\text{feature matrix } \mathbf{X}_{m,n+1}} \quad \underbrace{\begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}}_{\text{parameter vector } \boldsymbol{\theta}} \quad \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}}_{\text{label vector } \mathbf{y}_{m,1}}$$

Vectorized Form of the Gradient of MSE

$$\nabla \text{MSE}(h_{\theta}, D) = \frac{2}{m} \mathbf{X}^T (\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y})$$

$$\frac{2}{m} \mathbf{X}^T (\mathbf{X} \cdot \boldsymbol{\theta} - \mathbf{y}) = \mathbf{0}$$

$$\mathbf{X}^T \cdot \mathbf{X} \cdot \boldsymbol{\theta} = \mathbf{X}^T \cdot \mathbf{y}$$

$\boldsymbol{\theta} = \mathbf{X}^{\dagger} \cdot \mathbf{y}$, where $\mathbf{X}^{\dagger} = (\mathbf{X}^T \cdot \mathbf{X})^{-1} \cdot \mathbf{X}^T$ is the **pseudo-inverse** of \mathbf{X}

Additional Notes

- Luckily, there exists a closed-form solution to the convex optimization problem above (i.e., to minimize **MSE**).
- However, other choices of loss functions (even if convex) may need an **iterative** approach to get to a (local) minimum.
- For example, **gradient descent** use the gradient to **iteratively** converge to a local minimum (guaranteed for convex loss functions).
- We should use gradient descent as the size of our problem increases because computing the inverse of a large matrix is generally a very costly task ($O(n^3)$, where $n \times n$ is the size of the matrix).

Gradient Descent

- A generic iterative algorithm used to find an approximate solution to an optimization problem.
- It consists of 3 parameters:
 - $\theta_0 \in \mathbb{R}^n$ as the initial (random) guess
 - $\epsilon > 0$ stopping criterion
 - α step size

Gradient Descent: The Algorithm

1. $\theta^* = \theta_0$;
2. **while** $\|\nabla f\|_2 > \epsilon$: *// replace f with the function of interest, e.g., $\text{MSE}(h_\theta, \mathcal{D})$*
 $\theta^* = \theta^* - \underbrace{\alpha}_{\text{step size}} \nabla f(\theta^*)$; *// the direction of greatest decrease of f is opposite to the gradient vector*
3. **return** θ^* ;

Gradient Descent: The Choice of θ_0

- The starting point θ_0 can be chosen arbitrarily, though for non-convex function f the output of gradient descent can vary with its choice.
- For convex f , gradient descent will converge toward the **same point**, i.e., the global minimum, independently of the starting point.
- The choice of starting point can still affect the number of iterations until convergence.
- In practice, one should choose θ_0 according to our best guess as to where the global minimum is likely to be.

Gradient Descent: The Choice of ϵ

- The parameter ϵ determines the stopping criterion.
- Note that since $\epsilon > 0$, gradient descent generally does not halt at an **actual** local minimum, but rather at some kind of "approximate local minimum".
- Smaller values of ϵ mean more iterations before stopping but a higher-quality solution at termination.
- In practice, one tries various values of ϵ to achieve the right balance between computation time and solution quality.

Gradient Descent: The Choice of α

- The final parameter α , the "step size", is perhaps the most important.
- While gradient descent is flexible enough that different values of α can be used in different iterations, in practice one typically sets α in advance and stick to it over all iterations.
- The "best" value of α is typically chosen by experimentation. For example, we can run the entire gradient descent algorithm with a few different choices of α to see which run gives us the best results.


Overfitting (High Variance)

- Minimizing the loss function considering the whole dataset \mathcal{D} just limits the **in-sample error**.
- Our ultimate goal is to pick a hypothesis h^* which is able to **generalize** to unseen instances (i.e., minimize the **out-of-sample error**).
- Note that if we pick a hypothesis that just memorizes all the instances in \mathcal{D} , this will have a 0 **in-sample error** but this is **not** learning!
- **Overfitting:** h^* is not learning the true f but just mimic random noise
- **High Variance:** h^* does not generalize, as it strongly depends on \mathcal{D}


Underfitting (High Bias)

- Anyway, we don't want h^* to perform poorly on \mathcal{D} .
- **Underfitting:** h^* has a very high in-sample error
- **High Bias:** h^* does not take advantage of \mathcal{D}


Bias-Variance Tradeoff



High bias
(underfit)



"Just right"




High variance
(overfit)

Splitting the Dataset \mathcal{D}

- To prevent overfitting, split the dataset into (at least) two separate random samples:
 - **Training Set** ($\mathcal{D}_{\text{train}}$) used to learn h^* , such that $|\mathcal{D}_{\text{train}}| = 0.8 |\mathcal{D}|$
 - **Test or Held-out Set** ($\mathcal{D}_{\text{test}}$) used to evaluate h^* , such that $|\mathcal{D}_{\text{test}}| = |\mathcal{D}| - |\mathcal{D}_{\text{train}}|$
- Note that the (*offline* or *online*) metrics we use to evaluate h^* do not necessarily have to be the same of that we optimize at training time.

Splitting the Dataset \mathcal{D}



How Much Data Do We Need?

- In general, the more data we have the better we learn.
- Even if the datasets we operate on are not so big, we can still take advantage of ML.
- For example, we can use a technique called **k -fold cross validation** (e.g., $k = 10$) instead of just splitting \mathcal{D} in two datasets.


k -fold Cross Validation

- (Optional) Split the dataset in $\mathcal{D}_{\text{train}}$ and $\mathcal{D}_{\text{test}}$, as discussed above.
- Repeat for k iterations the following step:
 - Randomly select a proportion of $1/k$ instances from $\mathcal{D}_{\text{train}}$ which will be used for testing, whilst the remaining $(k - 1)/k$ proportion of the dataset is used for training.
 - Compute the error on the test proportion.
- Eventually, compute the **cross validation** error as the average of the individual errors obtained in each of the k runs.

4. Model Selection/Evaluation

Which Model Should I Use?

- Several learning models are designed to achieve the same tasks.
- Each learning model has its own set of parameters, a.k.a. **hyperparameters** (e.g., the number k of neighbors in k NN).



Validation Set

- It is crucial not to mix model training with hyperparameter sweeping.
- Another portion of the original dataset (**Validation Set**) should be reserved for this task.
- This can be also done within each step of a k -fold cross validation.

Example

Suppose we want to select which value of $k \in \{2, 5, 10\}$ of a k NN gives the best performance.

1. Train a separate model for each value of k on the training set (e.g., 70%)
2. Measure the error of each model on the validation set (e.g., 10%)
3. Select the model whose value of k gives the best performance on the validation set (e.g., $k = 5$)
4. Re-train **only** this model on the training + validation sets
5. Measure the performance on the test set (e.g., 20%)

Example Using Cross Validation

1. For each cross validation run:

- Train a separate model for each value of k on the training set portion of that run
- Measure the error of each model on the validation set portion of that run

2. Select the model whose value of k gives the best performance on the 10 averaged validation sets (e.g., $k = 5$)

3. Re-train **only** this model on the whole training set (e.g., 80%)

4. Measure the performance on the test set (e.g., 20%)